



Institut de Mathématiques de Toulouse, INSA Toulouse

Supervised Learning- Part II

Machine Learning for Data Science
CERFACS- May 2018

Béatrice Laurent-Philippe Besse- Aurélien Garivier

Methods studied in this course :

Part I

- Linear model, model selection, variable selection, Ridge regression, Lasso.
- Logistic regression
- Support Vector Machine

Part II

- Classification And Regression Trees (CART)
- Bagging, Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Outline

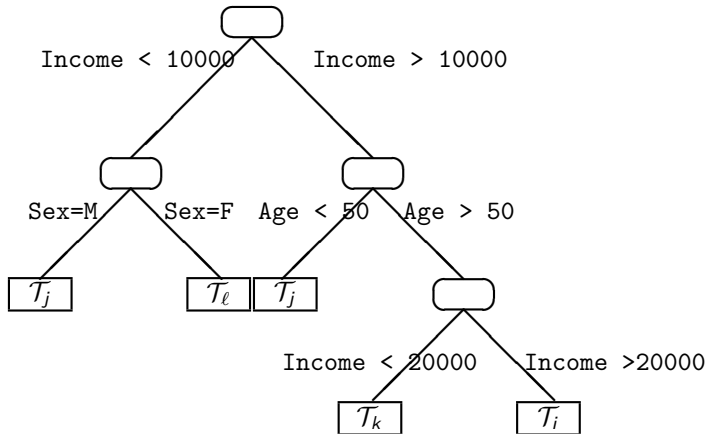
- Classification And Regression Trees (CART)
- Bagging, Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Classification And Regression Trees

Introduction

- Classification and regression trees (**CART**) : Breiman et al. (1984)
- X^j explanatory variables (quantitative or qualitative)
- Y qualitative with m modalities $\{\mathcal{T}_\ell; \ell = 1 \dots, m\}$: **classification tree**
- Y quantitative : **regression tree**
- **Objective** : construction of a binary **decision tree** easy to interpret
- No assumption on the model : non parametric procedure.



Example of binary classification tree



Definitions

- Determine an *iterative* sequence of *nodes*
- *Root* : initial note : the whole sample
- *Leaf* : Terminal node
- *Node* : choice of one *variable* and one *division* to proceed to a dichotomie
- *Division* : threshold value or group of modalities

Example of dyadic partition of the space



Rules

- We have to choose :
 - 1 Criterion for the “best” division among all *admissibles* ones (partition based on the values of one variable)
 - 2 Rules for a terminal node : *leaf*
 - 3 Rules to assign to a class \mathcal{T}_ℓ or one value for Y
- *Admissible* divisions : descendants $\neq \emptyset$
- X^j real or ordinal with c_j possible values : $(c_j - 1)$ possible divisions
- X^j nominal : $2^{(c_j-1)} - 1$ possible divisions
- Heterogeneity function D_κ of one node
 - 1 Null : a single modality for Y or Y is constant
 - 2 Maximal : all the modalities for Y or large variance

Division criterion

Optimal division

- Notations
 - κ : a node
 - κ_L and κ_R the two son nodes
- The algorithm retains the division which minimizes

$$D_{\kappa_L} + D_{\kappa_R}$$

- For each node κ in the construction of the tree :

$$\max_{\{ \text{Divisions of } X^j; j=1,p \}} D_{\kappa} - (D_{\kappa_L} + D_{\kappa_R})$$

Stopping rule and affectation

Leaf and affectation

- A **Node** is a **terminal** node or a **leaf**, if it is :
 - **Homogeneous**
 - **Number** of observations below some **threshold**
- **Affectation**
 - **Y quantitative**, the value is the **mean of the observations in the leaf**
 - **Y qualitative**, each leaf is affected to one class \mathcal{T}_ℓ of Y by considering the **conditional mode** :
 - the **mostly represented** class in the node
 - The **less costly** class if **cost for wrong classification** are given

Heterogeneity criterion in regression

Y quantitative : heterogeneity in regression

Heterogeneity of the node κ :

$$D_{\kappa} = \sum_{i \in \kappa} (y_i - \bar{y}_{\kappa})^2 = |\kappa| \frac{1}{|\kappa|} \sum_{i \in \kappa} (y_i - \bar{y}_{\kappa})^2$$

where $|\kappa|$ is the cardinality of the node κ

Minimize the intra-class variance

The son nodes κ_L and κ_R minimize :

$$\frac{1}{n} \sum_{i \in \kappa_L} (y_i - \bar{y}_{\kappa_L})^2 + \frac{1}{n} \sum_{i \in \kappa_R} (y_i - \bar{y}_{\kappa_R})^2.$$

Heterogeneity criterion in classification

Y qualitative : $Y \in \{\mathcal{T}_\ell, \ell = 1, \dots, m\}$.

Node κ .

- $n_\kappa^\ell = \text{Card} \{(X_i, Y_i), X_i \in \kappa, Y_i \in \mathcal{T}_\ell\}$.
- $n_\kappa = \text{Card} \{(X_i, Y_i), X_i \in \kappa\}$.
- p_κ^ℓ : probability that an observation is in class \mathcal{T}_ℓ given that it is in node κ .
- Estimated by $\frac{n_\kappa^\ell}{n_\kappa}$.

Heterogeneity criterion in classification

Y qualitative : heterogeneity in classification

Heterogeneity of node κ :

- **Shannon Entropy** with the notation $0 \log(0) = 0$
 p_{κ}^{ℓ} : proportion of the class \mathcal{T}_{ℓ} of Y in the node κ .

$$E_{\kappa} = - \sum_{\ell=1}^m p_{\kappa}^{\ell} \log(p_{\kappa}^{\ell})$$

Maximal in $(\frac{1}{m}, \dots, \frac{1}{m})$, minimal in $(1, 0, \dots, 0), \dots, (0, \dots, 0, 1)$.

$$D_{\kappa} = -|\kappa| \sum_{\ell=1}^m p_{\kappa}^{\ell} \log(p_{\kappa}^{\ell})$$

- **Gini concentration** : $D_{\kappa} = |\kappa| \sum_{\ell=1}^m p_{\kappa}^{\ell} (1 - p_{\kappa}^{\ell})$

Pruning and optimal tree

Pruning : notations

- We look for a **parsimonious** tree
- **Complexity** of a tree : K_A = numbers of leaves in A
- Adjustment error of A :

$$D(A) = \sum_{\kappa=1}^{K_A} D_{\kappa}$$

D_{κ} : heterogeneity of leaf κ

Sequence of embedded trees

- Adjustment error **penalized** by the **complexity** :

$$Crit_{\gamma}(A) = D(A) + \gamma \times K_A$$

- When $\gamma = 0$: A_{\max} (maximal tree) minimizes $Crit_{\gamma}(A)$
- When γ increases, the division of A_H , for which the improvement of D is smaller than γ , is cancelled ; **hence**
 - two leaves are gathered (**pruned**)
 - their father node becomes a **terminal** node
 - A_K becomes A_{K-1} .
- After **iteration** of this process, we get a sequence of trees :

$$A_{\max} \supset A_K \supset A_{K-1} \supset \cdots A_1$$

Breiman sub-sequence

- A_K is the sub tree of A_{\max} (maximal tree) obtained by pruning the nodes κ such that $D(\kappa) = D(\kappa_L) + D(\kappa_R)$.
- For each node in A_K , $D(\kappa) > D(\kappa_L) + D(\kappa_R)$ and $D(\kappa) > D(A_K^\kappa)$ where A_K^κ is the subtree of A_K from node κ .
- For γ small, $D(\kappa) + \gamma > D(A_K^\kappa) + \gamma|A_K^\kappa|$. This holds while $\gamma < (D(A_K^\kappa) - D(\kappa))/(|A_K^\kappa| - 1) = s(\kappa, A_K^\kappa)$ for all node κ of A_K .

$$\gamma_K = \inf_{\kappa \text{ node of } A_K} s(\kappa, A_K^\kappa)$$

- $Crit_{\gamma_K}(\kappa) = Crit_{\gamma_K}(A_K^\kappa)$ and the node κ becomes preferable to the subtree A_K^κ .
- $A_{K-1} = A_{\gamma_K}$ is the subtree obtained by pruning the branches from the nodes minimizing $s(\kappa, A_K^\kappa)$: this gives the second tree in the sub-sequence
- We iterate this process.

Optimal tree

Algorithm to select the optimal tree

- Maximal tree A_{\max}
- Imbedded sequence $A_{\max}, A_K \dots A_1$ associated with an increasing sequence of values $\gamma_K \leq \dots \leq \gamma_1$
- V-fold cross validation error :
 for $v = 1, \dots, V$ **do**
 - Estimation of the sequence of trees associated to (γ_K) with all the folds except v
 - Estimation of the error with the fold v .
- EndFor**
- Sequence of the mean of these errors for each value of $\gamma_K, \dots, \gamma_1$
- γ_{Opt} optimal value for the tuning parameter minimizing the mean of the errors
- Tree associated to γ_{Opt} in $A_K \dots A_1$

Advantages

- **Trees** are easy to interpret
- **Efficient algorithms** to find the pruned trees
- Tolerant to **missing data**


⇒ Success of CART for practical applications

Warnings

- **Variable selection** : the selected tree only depends on few explanatory variables, trees are often (wrongly) interpreted as a variable selection procedure
- **High instability** of the trees : not robust to the learning sample, curse of dimensionality ..
- **Prediction accuracy** of a tree is often poor compared to other procedures

⇒ **Aggregation of trees : bagging, random forests**

Example for Ozone data



Ozone : Classification tree pruned by cross-validation

Outline

- Classification And Regression Trees (CART)
- Bagging , Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Introduction

- Combination or **aggregation** of models (almost) without **overfitting**
 - **Bagging** is for **bootstrap aggregating** : Breiman, 1996
 - **Random forests** : Breiman, 2001
 - From **Boosting** (Freund et Shapiro, 1996) deterministic and **adaptative** to **extreme gradient boosting**
 - Allows to aggregate any modelisation method
 - **Efficient** methods : Fernandez-Delgado et al. (2014), *Kaggle*
-
- **Bagging** is appropriate for unstable algorithms, with small bias and high variance (CART)
 - **Boosting** is appropriate for algorithms with high bias and small variance

Bagging

Bootstrap aggregating

- Let Y a quantitative or qualitative variable to explain
- X^1, \dots, X^p the explanatory variables
- $f(\mathbf{x})$ a model function of $\mathbf{x} = \{x^1, \dots, x^p\} \in \mathbb{R}^p$
- $\mathbf{z} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ learning sample with distribution F and size n
 - $\hat{f}_{\mathbf{z}}$: predictor associated to the sample \mathbf{z} with $f(\cdot) = E_F(\hat{f}_{\mathbf{z}})$
 - B independent samples $\{\mathbf{z}_b\}_{b=1, B}$
 - Y quantitative : $\hat{f}_B(\cdot) = \frac{1}{B} \sum_{b=1}^B \hat{f}_{\mathbf{z}_b}(\cdot)$ (mean)
 - Y qualitative : $\hat{f}_B(\cdot) = \arg \max_j \text{card} \left\{ b \mid \hat{f}_{\mathbf{z}_b}(\cdot) = j \right\}$ (majority vote)
- **Principle** : Take the mean of independent predictions to reduce the variance
- B independent samples replaced by B bootstrap replications

Bagging : algorithm

- Let \mathbf{x}_0 to point where we want to predict and
- $\mathbf{z} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ a learning sample
- **For** $b = 1$ **to** B
 - Generate a bootstrap sample \mathbf{z}_b^* (with size $m_n \leq n$)
 - Estimate $\hat{f}_{\mathbf{z}_b}(\mathbf{x}_0)$ with the bootstrap sample
- Compute the mean estimation $\hat{f}_B(\mathbf{x}_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_{\mathbf{z}_b}(\mathbf{x}_0)$ or the result of the majority vote

- The B bootstrap samples are built on the same learning sample $\mathbf{z} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$:
 \implies the estimators $\hat{f}_{\mathbf{z}_b}(\mathbf{x}_0)$ are **not independent**
- Regression case : If $\text{Corr}(\hat{f}_{\mathbf{z}_b}(\mathbf{x}_0), \hat{f}_{\mathbf{z}_{b'}}(\mathbf{x}_0)) = \rho(\mathbf{x}_0)$,

$$\begin{aligned} E(\hat{f}_B(\mathbf{x}_0)) &= f(\mathbf{x}_0) \\ \text{Var}(\hat{f}_B(\mathbf{x}_0)) &= \rho(\mathbf{x}_0)\text{Var}(\hat{f}_b(\mathbf{x}_0)) + \frac{(1 - \rho(\mathbf{x}_0))}{B} \text{Var}(\hat{f}_b(\mathbf{x}_0)) \\ &\rightarrow \rho(\mathbf{x}_0)\text{Var}(\hat{f}_b(\mathbf{x}_0)) \text{ as } B \rightarrow +\infty \end{aligned}$$

\implies Importance to find **low correlated predictors** $(\hat{f}_b(\mathbf{x}_0))_{1 \leq b \leq B}$.
 \hookrightarrow **Random forests**

Bagging : practical use

- *Bootstrap out-of-bag* estimation of the prediction error : control of the **quality** and of **overfitting**
- **CART** to built a sequence of **binary trees**
- Three pruning **strategies** are possible :
 - 1 keep a **whole tree** for each of the samples
 - 2 tree with at most **q leaves**
 - 3 whole tree **pruned** by cross-validation
- **First strategy** compromise between computations and prediction accuracy : small **bias** for each tree and reduced **variance** by aggregation

Bagging : limitations

- Computation time and control of the error
- Storage of all the models to aggregate
- **Black box** model

Outline

- Classification And Regression Trees (CART)
- Bagging, Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Random forests

Random forests : principle

- **Random forests** means initially any aggregation method for classification or regression trees
- Now, it refers to the **Random input** method introduced by Breiman and Cutler (2005)
<http://www.stat.berkeley.edu/users/breiman/RandomForests/>
- Improvement of the **bagging** of binary trees
- Variance of B correlated variables : $\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$
- Add a **randomisation** to get more **independent** trees
- **Random** choice of variables
- **Interesting** in high dimension

Random forests : algorithm

- Let \mathbf{x}_0 the point where we want to predict,
 $\mathbf{z} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ a learning sample **for** $b = 1$ **to** B
 - Generate a bootstrap sample \mathbf{z}_b^*
 - Estimate a tree with randomization of the variables :
 - For each **node**, **random sample** of $m < p$ **predictors** to built the subdivision
- **Compute** the **mean** estimation $\hat{f}_B(\mathbf{x}_0) = \frac{1}{B} \sum_{b=1}^B \hat{f}_{\mathbf{z}_b}(\mathbf{x}_0)$
or the majority **vote**

Random forests : utilisation

- **Pruning** : tree with q leaves, or complete tree,
- **Random selection** of m predictors : $m = \sqrt{p}$ for classification, $\frac{p}{3}$ for regression.
- Small m increases the **variability** of each tree but reduces the correlation between them : Each **model** is less accurate but l'**aggregation** is performing
- Iterative computation of the **out-of-bag** error.
- Good accuracy, easily implementable, parallelisable but **not easy to interpret**

To help interpretation

- Index of importance for each variable
 - Mean Decrease Accuracy
 - Mean Decrease Gini

Other possible applications

- Proximity or similarity between observations
- Anomaly detection with IsolationForest
- Imputation of missing data with missForest
- Survival analysis with survival forest
- ...

Outline

- Classification And Regression Trees (CART)
- Bagging , Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Boosting principle

Boosting : principle

- Improve the performances of a **weak classifier** (Schapire, 1990 ; Freund and Schapire, 1996)
- **AdaBoost** (**Adaptative boosting**) prediction of a binary variable
- Reduction of the **variance** but also the **bias**
- Aggregation of a family of **recurrent** models : *Each **model** is an **adaptive** version of the previous one giving more **weight**, in the next estimate, to the **badly adjusted** observations*
- **Variants** : **type** of variable to predict (binary, k classes, real), **loss function**

Basic algorithm

AdaBoost

- AdaBoost is a boosting method to combine several binary classifiers f_1, \dots, f_k with values in $\{-1, 1\}$.
- $\mathbf{z} = \{(\mathbf{x}_1, y_1), \dots, (\mathbf{x}_n, y_n)\}$ a learning sample, $y_i \in \{-1, 1\}$.
- The aim is to minimize the empirical convexified risk for the exponential loss function over linear combination of the classifiers

$$\hat{f} = \underset{f \in \text{span}(f_1, \dots, f_k)}{\text{argmin}} \left\{ \frac{1}{n} \sum_{i=1}^n \exp(-y_i f(\mathbf{x}_i)) \right\}.$$

AdaBoost

- To approximate the solution, Adaboost computes a sequence of functions \hat{f}_m for $m = 0, \dots, M$ with

$$\begin{aligned}\hat{f}_0 &= 0 \\ \hat{f}_m &= \hat{f}_{m-1} + \beta_m f_{j_m}\end{aligned}$$

where (β_m, j_m) minimizes the empirical risk

$$\operatorname{argmin}_{\beta \in \mathbb{R}, j=1, \dots, p} \left\{ \frac{1}{n} \sum_{i=1}^n \exp(-y_i(\hat{f}_{m-1}(\mathbf{x}_i) + \beta f_j)) \right\}.$$

- The final classification rule is given by

$$\hat{f} = \operatorname{sign}(\hat{f}_M).$$

AdaBoost

- We denote

$$w_i^{(m)} = \frac{1}{n} \exp(-y_i \hat{f}_{m-1}(\mathbf{x}_i))$$

- We assume that for all $j = 1, \dots, k$,

$$\widehat{\mathcal{E}}_m(j) = \frac{\sum_{i=1}^n w_i^{(m)} \mathbb{1}_{f_j(\mathbf{x}_i) \neq y_i}}{\sum_{i=1}^n w_i^{(m)}} \in]0, 1[.$$

- Then, we have :

$$j_m = \underset{j=1, \dots, k}{\operatorname{argmin}} \widehat{\mathcal{E}}_m(j),$$

and

$$\beta_m = \frac{1}{2} \log \left(\frac{1 - \widehat{\mathcal{E}}_m(j)}{\widehat{\mathcal{E}}_m(j)} \right).$$

AdaBoost

AdaBoost algorithm

- $w_i^{(1)} = 1/n$ for $i = 1, \dots, n$.
- For $m = 1, \dots, M$

$$j_m = \underset{j=1,\dots,p}{\operatorname{argmin}} \widehat{\mathcal{E}}_m(j)$$

$$\beta_m = \frac{1}{2} \log \left(\frac{1 - \widehat{\mathcal{E}}_m(j)}{\widehat{\mathcal{E}}_m(j)} \right)$$

$$\begin{aligned} w_i^{(m+1)} &= w_i^{(m)} \exp(-y_i \beta_m f_{j_m}(\mathbf{x}_i)) \text{ for } i = 1, \dots, n \\ &= \frac{1}{n} \exp(-y_i \hat{f}_m(\mathbf{x}_i)) \end{aligned}$$

- $\hat{f}_M(x) = \sum_{m=1}^M \beta_m f_{j_m}(x)$.
- $\hat{f} = \operatorname{sign}(\hat{f}_M)$.

Boosting : utilisation

- Tree as basic model
- Recommendation : q between 4 and 8
- Empirically , the prediction error can still decrease while the training error is equal to zero
- Warning : noisy data (error in the label), source of drift or overfitting
- Boosting reduces the variance like the bagging but also the bias

Gradient Boosting Models

GBM : Principle 1

- Gradient Boosting Models (Friedman, 2002-2009)
- In the case of a **differentiable** loss function
- **Principle** :
 - Construct a **sequence** of models in such a way that at each **step**, each **model** added to the **linear combination**, appears as a **step** towards a **better solution**
 - This **step** is done in the direction of the **gradient** of the **loss function** approached by a **regression tree**

GBM : Principle 2

- Previous **Adaptative model** (AdaBoost) :

$$f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \beta_m f_{j_m}(\mathbf{x})$$

- Transformed in a **gradient descent**

$$f_m = f_{m-1} - \gamma_m \sum_{i=1}^n \nabla_{f_{m-1}(\mathbf{x}_i)} l(y_i, f_{m-1}(\mathbf{x}_i)).$$

- Search for a **best step in the descent** γ :

$$\min_{\gamma} \sum_{i=1}^n \left[l \left(y_i, f_{m-1}(\mathbf{x}_i) - \gamma \frac{\partial l(y_i, f_{m-1}(\mathbf{x}_i))}{\partial f_{m-1}(\mathbf{x}_i)} \right) \right].$$

GBM in regression : algorithm

- Let \mathbf{x}_0 the point where we want to predict
- Initialize $\hat{f}_0 = \arg \min_{\gamma} \sum_{i=1}^n l(y_i, \gamma)$
- **For** $m = 1$ **to** M
 - Compute $r_{mi} = - \left[\frac{\delta l(y_i, f(\mathbf{x}_i))}{\delta f(\mathbf{x}_i)} \right]_{f=f_{m-1}} ; \quad i = 1, \dots, n$
 - Adjust a regression tree δ_m to (\mathbf{x}_i, r_{mi})
 - Calculate $\gamma_m = \arg \min_{\gamma} \sum_{i=1}^n l(y_i, f_{m-1}(\mathbf{x}_i) + \gamma \delta_m(\mathbf{x}_i))$
 - Update : $\hat{f}_m(\mathbf{x}) = \hat{f}_{m-1}(\mathbf{x}) + \gamma_m \delta_m(\mathbf{x})$
- **Result** : $\hat{f}_M(\mathbf{x}_0)$

GBM : use with R

- **Discrimination** : as many probabilities as classes
- **Shrinkage** coefficient

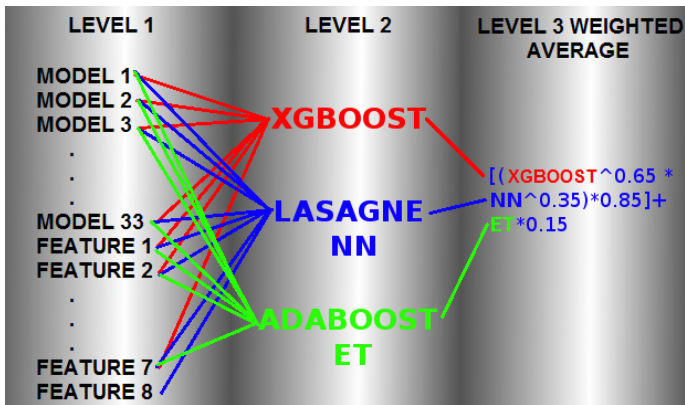
$$\hat{f}_m(\mathbf{x}) = \hat{f}_{m-1}(\mathbf{x}) + \nu \sum_{j=1}^{J_m} \gamma_{jm} \mathbf{1}\{\mathbf{x} \in R_{jm}\}$$

- **Trade-off** between shrinkage and number of iterations M .
- Maximum **depth** of the trees

Extreme gradient boosting

XGBoost : motivation

- **Algorithm** introduced by Chen et Guestrin (2016)
- Additional **Penalization** to control overfitting
- **Problem** : number of parameters to optimize
- Implementation **tips** for parallelisation
- **Environments** : R (caret), Python, Julia, GPU, *Amazon Web Service*, Spark...
- **Winning solutions** for *Kaggle* competitions



Kaggle : Identify people who have a high degree of Psychopathy based on Twitter usage

XGBoost : penalization

- Loss function L penalized version of l

$$L(f) = \sum_{i=1}^n l(f(\mathbf{x}_i), y_i) + \sum_{m=1}^M \Omega(\delta_m)$$

$$\Omega(\delta) = \alpha |\delta| + \beta \|\mathbf{w}\|^2$$

- $|\delta|$ number of leaves in the tree δ
- \mathbf{w} vector of values assigned to each leaf
- Ω Mix of l_1 and l_2 penalization

XGBoost : tips

- **Gradient Approximation** by Taylor expansion : summations and parallelization
- **Complexity** of subdivisions in the trees : quantiles of the distributions
- Tolerant algorithm for **Missing data** : gradient calculated on present values
- **Index of importance** of the variables

XGBoost : additional parameters

- alpha** Lasso (l_1) penalization of the complexity of the regression tree to estimate the gradient
- lambda** Ridge (l_2) penalization
- gamma** Minimal reduction of the loss to accept a division
- tree_method** greedy algorithm for searching for divisions or simplification (quantiles or grouping of classes)
- sketch_eps** control of the number of classes
- scale_pos_weight** to take into account for unbalanced classes
- Other** Optimisation parameters of the performances

XGBoost : Optimisation strategy

- Principle :
- **Default values** and optimize the parameters by decreasing order of the supposed influence
- A **strategy** among others :
 - 1 Number of trees
 - 2 Maximal depth vs. minimal number observations per leaf
 - 3 Minimal reduction of the loss
 - 4 Sampling rate vs. number of variables used
 - 5 Number of trees vs. shrinkage
- **Need** : computing power (GPU)

Outline

- Classification And Regression Trees (CART)
- Bagging, Random Forests
- Boosting
- Neural networks, Introduction to deep learning

Introduction

- Deep learning is a set of learning methods attempting to model data with complex architectures combining different non-linear transformations.
- The elementary bricks of deep learning are the neural networks, that are combined to form the deep neural networks.

There exist several types of architectures for neural networks :

- The multilayer perceptrons, that are the oldest and simplest ones
- The Convolutional Neural Networks (CNN), particularly adapted for image processing
- The recurrent neural networks, used for sequential data such as text or times series.

Introduction

- Deep learning architectures are based on **deep cascade of layers**.
- They need clever **stochastic optimization algorithms**, and **initialization**, and also a clever choice of the **structure**.
- They lead to very **impressive results**, although very **few theoretical foundations** are available till now.
- These techniques have enabled **significant progress** in the fields of **sound and image processing**, including facial recognition, speech recognition, computer vision, automated language processing, text classification (for example spam recognition).

Neural networks

- An **artificial neural network** is non linear with respect to its parameters θ that associates to an entry x an output $y = f(x, \theta)$.
- The neural networks can be use for **regression or classification**.
- The parameters θ are estimated from a learning sample.
- The function to minimize is not convex, leading to local minimizers.
- The success of the method came from a **universal approximation theorem** due to Cybenko (1989) and Hornik (1991).
- Le Cun (1986) proposed an efficient way to compute the gradient of a neural network, called **backpropagation of the gradient**, that allows to obtain a local minimizer of the quadratic criterion easily.

Artificial Neuron

Artificial neuron

- a function f_j of the input $x = (x_1, \dots, x_d)$
- weighted by a vector of **connection weights** $w_j = (w_{j,1}, \dots, w_{j,d})$,
- completed by a **neuron bias** b_j ,
- and associated to an **activation function** ϕ :

$$y_j = f_j(x) = \phi(\langle w_j, x \rangle + b_j).$$

Activation functions

Activation functions

- The identity function $\phi(x) = x$
- The sigmoid function (or logistic) $\phi(x) = \frac{1}{1+e^{-x}}$
- The hyperbolic tangent function ("tanh")

$$\phi(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$

- The hard threshold function $\phi_{\beta}(x) = \mathbb{1}_{x \geq \beta}$
- The Rectified Linear Unit (ReLU) activation function
 $\phi(x) = \max(0, x)$

Activation functions





FIGURE: Activation functions

Artificial Neuron

Schematic representation of an **artificial neuron** where $\Sigma = \langle w_j, x \rangle + b_j$.



- Historically, the **sigmoid** was the mostly used activation function since it is differentiable and allows to keep values in the interval $[0, 1]$.
- Nevertheless, it is problematic since its gradient is very close to 0 when $|x|$ is not close to 0.




FIGURE: Sigmoid function (in black) and its derivatives (in red)

Activation functions

- With neural networks with a high number of layers, this causes troubles for the **backpropagation algorithm** to estimate the parameters.
- This is why the sigmoid function was supplanted by the **rectified linear function (ReLU)**.
- This function is not differentiable in 0 but the probability to have an entry equal to 0 is generally null.
- The ReLU function and its derivative are equal to 0 for negative values, hence it is advised
 - to add a small positive bias to ensure that each unit is active.
 - to consider a **variations of the ReLU function** such as

$$\phi(x) = \max(x, 0) + \alpha \min(x, 0)$$

where α is either a fixed parameter set to a small positive value, or a parameter to estimate.

Multilayer perceptron

- A **multilayer perceptron** (or neural network) is a structure composed by **several hidden layers of neurons** where the **output of a neuron of a layer becomes the input of a neuron of the next layer**.
- On last layer, called output layer, we may apply a **different activation function** as for the hidden layers depending on the type of problems we have at hand : **regression or classification**.




FIGURE: A basic neural network.

Multilayers perceptrons

- The parameters of the architecture are the **number of hidden layers and of neurons in each layer**.
- The **activation functions** are also to choose by the user. On the **output layer**, we apply **no activation function in the case of regression**.
- For **binary classification**, the output gives a prediction of $\mathbb{P}(Y = 1/X)$ since this value is in $[0, 1]$, **the sigmoid activation function** is generally considered.
- For **multi-class classification**, the output layer contains one neuron per class i , giving a prediction of $\mathbb{P}(Y = i/X)$. The sum of all these values has to be equal to 1.
- The multidimensional function **softmax** is generally used

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}.$$

Mathematical formulation

Multilayer perceptron with L hidden layers :

- We set $h^{(0)}(x) = x$.

For $k = 1, \dots, L$ (hidden layers),

$$\begin{aligned}a^{(k)}(x) &= b^{(k)} + W^{(k)}h^{(k-1)}(x) \\h^{(k)}(x) &= \phi(a^{(k)}(x))\end{aligned}$$

For $k = L + 1$ (output layer),

$$\begin{aligned}a^{(L+1)}(x) &= b^{(L+1)} + W^{(L+1)}h^{(L)}(x) \\h^{(L+1)}(x) &= \psi(a^{(L+1)}(x)) := f(x, \theta).\end{aligned}$$

where ϕ is the activation function and ψ is the output layer activation function

- At each step, $W^{(k)}$ is a matrix with number of rows the number of neurons in the layer k and number of columns the number of neurons in the layer $k - 1$.

Estimation of the parameters

- Once the **architecture of the network** has been chosen, the **parameters** (the weights w_j and biases b_j) have to be **estimated from a learning sample**.
- As usual, the estimation is obtained by **minimizing a loss function with a gradient descent algorithm**.
- We first have to choose the loss function.
- It is classical to estimate the parameters by minimizing the loss function which is the opposite of the log likelihood.
- Denoting θ the vector of parameters to estimate, we consider the expected loss function

$$L(\theta) = -\mathbb{E}_{(X,Y) \sim P}(\log(p_{\theta}(Y/X))).$$

Loss functions

- If the model is **Gaussian**, namely if $P_{\theta}(Y/X = x) \sim \mathcal{N}(f(x, \theta), I)$, maximizing the likelihood is equivalent to minimize the quadratic loss

$$L(\theta) = \mathbb{E}_{(X,Y) \sim P} (\|Y - f(X, \theta)\|^2).$$

- For **binary classification**, with $Y \in \{0, 1\}$, maximizing the log likelihood corresponds to the **minimization of the cross-entropy**. Setting $f(X, \theta) = P_{\theta}(Y = 1/X)$,

$$L(\theta) = -\mathbb{E}_{(X,Y) \sim P} [Y \log(f(X, \theta)) + (1 - Y) \log(1 - f(X, \theta))].$$

- This loss function is well adapted with the sigmoid activation function since the use of the logarithm avoids to have too small values for the gradient.

Loss functions

- For a **multi-class classification** problem, we consider a generalization of the previous loss function to k classes

$$L(\theta) = -\mathbb{E}_{(X,Y) \sim P} \left[\sum_{j=1}^k \mathbb{1}_{Y=j} \log p_{\theta}(Y = j/X) \right].$$

- Ideally we would like to minimize the **classification error**, but it is not smooth, this is why we consider the **cross-entropy** (or eventually a convex surrogate).

Penalized empirical risk

- The **expected loss** can be written as

$$L(\theta) = \mathbb{E}_{(X,Y) \sim P} [\ell(f(X, \theta), Y)]$$

and it is associated to a loss function ℓ .

- In order to estimate the parameters θ , we use a training sample $(X_i, Y_i)_{1 \leq i \leq n}$ and we **minimize the empirical loss**

$$\tilde{L}_n(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(X_i, \theta), Y_i)$$

eventually we add a **regularization term**.

- This leads to minimize the **penalized empirical risk**

$$L_n(\theta) = \frac{1}{n} \sum_{i=1}^n \ell(f(X_i, \theta), Y_i) + \lambda \Omega(\theta).$$

Penalized empirical risk

- We can consider \mathbb{L}^2 regularization.

$$\begin{aligned}\Omega(\theta) &= \sum_k \sum_i \sum_j (W_{i,j}^{(k)})^2 \\ &= \sum_k \|W^{(k)}\|_F^2\end{aligned}$$

where $\|W\|_F$ denotes the Frobenius norm of the matrix W .

- Note that only the weights are penalized, the biases are not penalized.
- It is easy to compute the gradient of $\Omega(\theta)$:

$$\nabla_{W^{(k)}} \Omega(\theta) = 2W^{(k)}.$$

- One can also consider \mathbb{L}^1 regularization, leading to **parcimonious solutions** :

$$\Omega(\theta) = \sum_k \sum_i \sum_j |W_{i,j}^{(k)}|.$$

Penalized empirical risk

- In order to minimize the criterion $L_n(\theta)$, a **stochastic gradient descent algorithm** is used.
- In order to compute the gradient, a clever method, called **Backpropagation algorithm** is considered.
- It has been introduced by Rumelhart et al. (1988), it is still **crucial** for deep learning.

The stochastic gradient descent algorithm

- Initialization of $\theta = (W^{(1)}, b^{(1)}, \dots, W^{(L+1)}, b^{(L+1)})$.
- **For** $j = 1, \dots, N$ iterations :
 - At step j :

$$\theta = \theta - \varepsilon \frac{1}{m} \sum_{i \in B} [\nabla_{\theta} \ell(f(X_i, \theta), Y_i) + \lambda \nabla_{\theta} \Omega(\theta)] .$$

The stochastic gradient descent algorithm

- In the previous algorithm, we do not compute the gradient for the loss function at each step of the algorithm but only on a subset B of cardinality m (called a **batch**).
- This is what is classically done for **big data sets (and for deep learning)** or for **sequential data**.
- B is taken **at random without replacement**.
- An iteration over all the training examples is called an **epoch**.
- The **numbers of epochs** to consider is a parameter of the deep learning algorithms.
- The total number of iterations equals the number of epochs times the sample size n divided by m , the size of a batch.
- This procedure is called **batch learning**, sometimes, one also takes batches of size 1, reduced to a single training example (X_i, Y_i) .

Backpropagation algorithm for regression

- We consider the **regression case** and explain how to compute the gradient of the empirical **quadratic loss** by the **Backpropagation algorithm**.
- To simplify, we do not consider here the penalization term, that can easily be added.
- Assuming that the output of the multilayer perceptron is of size K , and using the previous notations, the **empirical quadratic loss** is proportional to

$$\sum_{i=1}^n R_i(\theta)$$

with

$$R_i(\theta) = \sum_{k=1}^K (Y_{i,k} - f_k(X_i, \theta))^2.$$

Backpropagation algorithm for regression

- In a regression model, the output activation function ψ is generally the identity function, to be more general, we assume that

$$\psi(a_1, \dots, a_K) = (g_1(a_1), \dots, g_K(a_K))$$

where g_1, \dots, g_K are functions from \mathbb{R} to \mathbb{R} .

- Computation of the partial derivatives of R_i with respect to the weights of the output layer.
 - Recalling that

$$a^{(L+1)}(x) = b^{(L+1)} + W^{(L+1)} h^{(L)}(x),$$

we get

$$\frac{\partial R_i}{\partial W_{k,m}^{(L+1)}} = -2(Y_{i,k} - f_k(X_i, \theta)) g'_k(a_k^{(L+1)}(X_i)) h_m^{(L)}(X_i).$$

Backpropagation algorithm for regression

- Computation of the partial derivatives of R_i with respect to the weights of the previous layer.
 - Differentiating now with respect to the **weights of the previous layer**

$$\frac{\partial R_i}{\partial W_{m,l}^{(L)}} = -2 \sum_{k=1}^K (Y_{i,k} - f_k(X_i, \theta)) g'_k(a_k^{(L+1)}(X_i)) \frac{\partial a_k^{(L+1)}(X_i)}{\partial W_{m,l}^{(L)}}.$$

with

$$\begin{aligned} a_k^{(L+1)}(x) &= \sum_j W_{k,j}^{(L+1)} h_j^{(L)}(x), \\ h_j^{(L)}(x) &= \phi \left(b_j^{(L)} + \langle W_j^{(L)}, h^{(L-1)}(x) \rangle \right). \end{aligned}$$

- This leads to

$$\frac{\partial a_k^{(L+1)}(x)}{\partial W_{m,l}^{(L)}} = W_{k,m}^{(L+1)} \phi' \left(b_m^{(L)} + \langle W_m^{(L)}, h^{(L-1)}(x) \rangle \right) h_l^{(L-1)}(x).$$

Backpropagation algorithm for regression

- Let us introduce the notations

$$\begin{aligned}\delta_{k,i} &= -2(Y_{i,k} - f_k(X_i, \theta))g'_k(a_k^{(L+1)}(X_i)) \\ s_{m,i} &= \phi' \left(a_m^{(L)}(X_i) \right) \sum_{k=1}^K W_{k,m}^{(L+1)} \delta_{k,i}.\end{aligned}$$

- Then we have

$$\frac{\partial R_i}{\partial W_{k,m}^{(L+1)}} = \delta_{k,i} h_m^{(L)}(X_i) \quad (1)$$

$$\frac{\partial R_i}{\partial W_{m,l}^{(L)}} = s_{m,i} h_l^{(L-1)}(X_i), \quad (2)$$

known as the **backpropagation equations**.

Backpropagation algorithm for regression

- The values of the gradient are used to **update the parameters** in the gradient descent algorithm.
- At step $r + 1$, we have :

$$W_{k,m}^{(L+1,r+1)} = W_{k,m}^{(L+1,r)} - \varepsilon_r \sum_{i \in B} \frac{\partial R_i}{\partial W_{k,m}^{(L+1,r)}}$$
$$W_{m,l}^{(L,r+1)} = W_{m,l}^{(L,r)} - \varepsilon_r \sum_{i \in B} \frac{\partial R_i}{\partial W_{m,l}^{(L,r)}}$$

where B is a batch and $\varepsilon_r > 0$ is the **learning rate** that satisfies $\varepsilon_r \rightarrow 0$, $\sum_r \varepsilon_r = \infty$, $\sum_r \varepsilon_r^2 < \infty$, for example $\varepsilon_r = 1/r$.

Backpropagation algorithm for regression

- We use the **Backpropagation equations** to compute the gradient by a **two pass algorithm**.
- In the *forward pass*, we fix the value of the current weights $\theta^{(r)} = (W^{(1,r)}, b^{(1,r)}, \dots, W^{(L+1,r)}, b^{(L+1,r)})$, and we compute the predicted values $f(X_i, \theta^{(r)})$ and all the intermediate values $(a^{(k)}(X_i), h^{(k)}(X_i) = \phi(a^{(k)}(X_i)))_{1 \leq k \leq L+1}$ that are stored.
- Using these values, we compute during the *backward pass* the quantities $\delta_{k,i}$ and $s_{m,i}$ and the **partial derivatives given in Equations 1 and 2**.
- We have computed the partial derivatives of R_i only with respect to the weights of the output layer and the previous ones, but we can go on to compute the partial derivatives of R_i with respect to the weights of the previous hidden layers.
- In the back propagation algorithm, each hidden layer gives and receives informations from the neurons it is connected with.
- Hence, the algorithm is **adapted for parallel computations**.

Initialization

- The input data have to be **normalized** to have approximately the **same range**.
- The **biases** can be initialized to 0. The **weights** cannot be initialized to 0 since for the tanh activation function, the derivative at 0 is 0, this is a saddle point.
- They also cannot be initialized with the same values, otherwise, all the neurons of a hidden layer would have the same behaviour.
- We generally **initialize the weights at random** : the values $W_{i,j}^{(k)}$ are i.i.d. Uniform on $[-c, c]$ with possibly $c = \frac{\sqrt{6}}{N_k + N_{k-1}}$ where N_k is the size of the hidden layer k . We also sometimes initialize the weights with a normal distribution $\mathcal{N}(0, 0.01)$ (see Gloriot and Bengio, 2010).

Optimization algorithms

- Many algorithms can be used to minimize the loss function with hyperparameters to calibrate.
- The **Stochastic Gradient Descent (SGD)** algorithm :

$$\theta_i^{new} = \theta_i^{old} - \varepsilon \frac{\partial L}{\partial \theta_i}(\theta_i^{old}),$$

where ε is the *learning rate* , and its calibration is **very important** for the convergence of the algorithm.

- If it is **too small**, the convergence is **very slow** and the optimization can be blocked on a local minimum.
- If the learning rate is **too large**, the network will **oscillate** around an optimum without stabilizing and converging.
- A classical way to proceed is to **adapt the learning rate during the training** : recommended to begin with a "large " value of ϵ , and reduce this value during the **successive iterations** : The **observation of the evolution of the loss function** can give indications on the way to proceed.

Optimization algorithms

- Another stopping rule, called *early stopping* is also used : we stop learning when the loss function for a validation sample stops to decrease.
- *Batch learning* is used for computational reasons, backpropagation algorithms need to *store intermediate values* : for *big data sets*, such as millions of images, this is not feasible, all the more that the deep networks have *millions of parameters to calibrate*.
- The *batch size m* is also a parameter to *calibrate*. Small batches generally lead to better generalization properties.
- The particular case of batches of size 1 is called *On-line Gradient Descent*. The disadvantage of this procedure is the very long computation time.

SGD algorithm

Summary of Stochastic Gradient Descent algorithm

- Fix the parameters ϵ : learning rate, m : batch size, nb : number of epochs.
- For $l = 1$ to nb epochs
- For $i = 1$ to n/m ,
 - Take a random batch of size m without replacement in the learning sample : $(X_i, Y_i)_{i \in B_l}$
 - Compute the gradients with the backpropagation algorithm

$$\tilde{\nabla}_{\theta} = \frac{1}{m} \sum_{i \in B_l} \nabla_{\theta} \ell(f(X_i, \theta), Y_i).$$

- Update the parameters

$$\theta^{new} = \theta^{old} - \epsilon \tilde{\nabla}_{\theta}.$$

SGD algorithm


- Since the choice of the learning rate is delicate, variations of the algorithm that are less sensitive to this parameter have been proposed.
- **Nesterov accelerated gradient** : Nesterov (1983) and Sutskever et al. (2013)
- **RMSPprop** algorithm, due to Hinton (2012) .

Regularization


- We have already mentioned \mathbb{L}^2 or \mathbb{L}^1 penalization ; we have also mentioned **early stopping**.
- For deep learning, the mostly used method is the **dropout** introduced by Hinton et al. (2012).
- With a certain **probability p** , and independently of the others, each unit of the network is **set to 0**.
- It is classical to set p to **0.5** for units in the **hidden layers**, and to **0.2** for the **entry layer**.
- The **computational cost is weak** since we just have to set to 0 some weights with probability p .

Dropout

- This method **improves significantly** the generalization properties of deep neural networks and is now the **most popular regularization method** in this context.
- The disadvantage is that **training is much slower** (it needs to increase the number of epochs).



(a) Standard Neural Net



(b) After applying dropout.

FIGURE: Dropout

Convolutional neural networks

- For some types of data, especially for **images**, multilayer perceptrons are **not well adapted**.
- They are defined for **vectors**. By transforming the images into vectors, we lose **spatial informations**, such as forms.
- The **convolutional neural networks (CNN)** introduced by LeCun (1998) have revolutionized image processing.
- CNN act directly on **matrices**, or even on **tensors** for images with three RGB color channels.

Convolutional neural networks

- CNN are now widely used for image classification, image segmentation, object recognition, face recognition ..



FIGURE: Image annotation

Convolutional neural networks




FIGURE: Image Segmentation.

Layers in a CNN

- A Convolutional Neural Network is composed by several kinds of layers, that are described in this section :
 - convolutional layers
 - pooling layers
 - fully connected layers

Convolution layer

- For 2-dimensional signals such as images, we consider the **2D-convolutions** : $(K * I)(i, j) = \sum_{m, n} K(m, n)I(i + n, j + m)$.
- K is a **convolution kernel** applied to a 2D signal (or image) I .




FIGURE: 2D convolution

Convolution layer

- The principle of **2D convolution** is to drag a convolution kernel on the image.
- At each position, we get the **convolution between the kernel and the part of the image that is currently treated**.
- Then, the kernel moves by a number s of pixels, s is called the **stride**.
- When the stride is small, we get redundant information.
- Sometimes, we also add a **zero padding**, which is a margin of size p containing zero values around the image in order to control the **size of the output**.

Convolution layer

- Assume that we apply C_0 kernels, each of size $k \times k$ on an image.
- If the size of the input image is $W_i \times H_i \times C_i$ (W_i denotes the width, H_i the height, and C_i the number of channels, typically $C_i = 3$), the volume of the output is $W_0 \times H_0 \times C_0$, where C_0 corresponds to the number of kernels that we consider, and

$$W_0 = \frac{W_i - k + 2p}{s} + 1$$

$$H_0 = \frac{H_i - k + 2p}{s} + 1.$$





FIGURE: 2D convolution - Units corresponding to the same position but at various depths : each unit applies a different kernel on the same patch of the

Convolution layer



- Convolution (3-dim dot product) image and filter
- Stack filter in one layer (See blue and green output, called **channel**)

- If the image has **3 channels** and if K_l ($l = 1, \dots, C_0$) denote $5 \times 5 \times 3$ kernels, the convolution with the image I with the kernel K_l corresponds to the formula :

$$K_l * I(i, j) = \sum_{c=0}^2 \sum_{n=0}^4 \sum_{m=0}^4 K_l(n, m, c) I(i + n - 2, i + m - 2, c).$$

Convolution layer

- For images with C^i channels, the shape of the kernel is (k, k, C^i, C^0) where C^0 is the number of **output channels**.
- The convolution operations are combined with an **activation function** ϕ (ReLU in general) : if we consider a kernel K of size $k \times k$, if \mathbf{x} is a $k \times k$ patch of the image, the activation is obtained by sliding the $k \times k$ window and computing $\mathbf{z}(\mathbf{x}) = \phi(K * \mathbf{x} + b)$, where b is a bias.
- **CNN learn the filters** (or kernels) that are the **most useful** for the task that we have to do (such as **classification**). Several convolution layers are considered : the output of a convolution becomes the input of the next one.

Pooling layer

- CNN also have *pooling layers*, which allow to reduce the dimension, also referred as *subsampling*, by taking the *mean or the maximum* on patches of the image (*mean-pooling or max-pooling*).
- Like the convolutional layers, pooling layers acts on *small patches of the image*.
- If we consider 2×2 patches, over which we take the maximum value to define the output layer, with a stride 2, we *divide by 4* the size of the image.

Pooling layer

- Another advantage of the pooling is that it makes the network **less sensitive to small translations** of the input images.




FIGURE: Maxpooling and effect on the dimension

Fully connected layers

- After several convolution and pooling layers, the CNN generally ends with several *fully connected layers*.
- The tensor that we have at the output of these layers is transformed into a vector and then we add several *perceptron layers*.


Architectures

- We have described the different types of layers composing a CNN.
- We now present how this layers are combined to form the architecture of the network.
- Choosing an architecture is very complex and this is more engineering than an exact science.
- It is therefore important to study the architectures that have proved to be effective and to draw inspiration from these famous examples.
- In the most classical CNN, we chain several times a convolution layer followed by a pooling layer and we add at the end fully connected layers.

Architectures

The **LeNet** network, proposed by the inventor of the CNN, **Yann LeCun (1998)** is of this type. This network was devoted to **digit recognition**. It is composed only on few layers and few filters, due to the computer limitations at that time.

Putting It All Together!



Lenet-5 (Lecun-98), Convolutional Neural Network for digits recognition

• Lecun 1998

• 51

FIGURE: Architecture of the network Le Net. LeCun, Y., Bottou, L., Bengio, Y. and Haffner, P. (1998)

Architectures

- With the appearance of GPU (Graphical Processor Unit) cards, much more complex architectures for CNN have been proposed, like the network **AlexNet** (Krizhevsky (2012)).
- **AlexNet** won the ImageNet competition devoted to the classification of one million of color images (224×224) onto 1000 classes.
- AlexNet is composed of 5 convolution layers, 3 max-pooling 2×2 layers and fully connected layers.

Architectures




FIGURE: Architecture of the network AlexNet. Krizhevsky, A. et al (2012)

Architectures

Alexnet architecture

Input	227 * 227 * 3			
Conv 1	55*55*96	96	11 *11	filters at stride 4, pad 0
Max Pool 1	27*27*96		3 *3	filters at stride 2
Conv 2	27*27*256	256	5*5	filters at stride 1, pad 2
Max Pool 2	13*13*256		3 *3	filters at stride 2
Conv 3	13*13*384	384	3*3	filters at stride 1, pad 1
Conv 4	13*13*384	384	3*3	filters at stride 1, pad 1
Conv 5	13*13*256	256	3*3	filters at stride 1, pad 1
Max Pool 3	6*6*256		3 *3	filters at stride 2
FC1	4096	4096	neurons	
FC2	4096	4096	neurons	
FC3	1000	1000	neurons	(softmax logits)

Architectures

We present another example.




FIGURE: Simonyan, K. and Zisserman, A. Very deep convolutional networks for large-scale image recognition (2014).

Architectures

The network that won the competition in 2014 is the network **GoogLeNet**(Szegedy et al. (2016)), which is a new kind of CNN, not only composed on successive convolution and pooling layers, but also on new modules called *Inception*, which are some kind of network in the network.

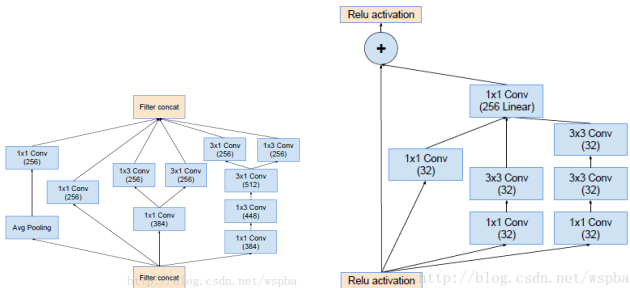


FIGURE: Inception modules, Szegedy et al. (2016)

Architectures

- The most recent innovations concern the **ResNet networks** (see Szegedy 2016).
- The originality of the ResNets is to add a **connection linking the input of a layer** (or a set of layers) with its output.
- In order to reduce the number of parameters, the ResNets do not have fully connected layers.
- GoogleNet and ResNet are **much deeper** than the previous CNN, but contain **much less parameters**.
- They are nevertheless **much costly in memory** than more classical CNN such as VGG or AlexNet.

Architectures

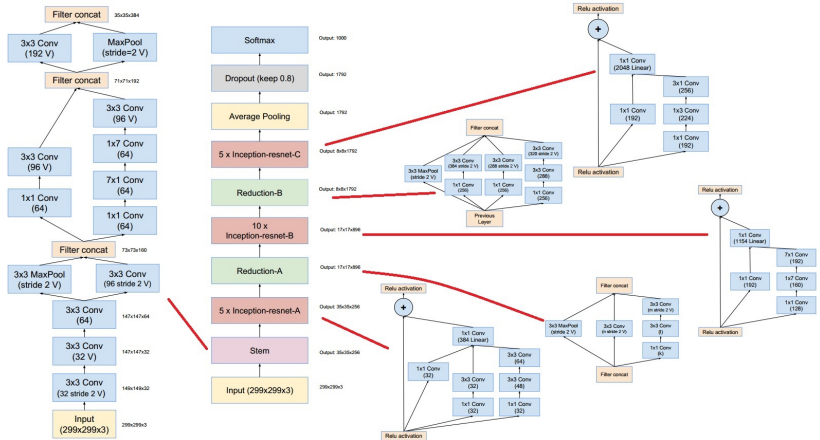


FIGURE: Inception-v4, Inception-resnet (Szegedy, C. et al. , 2016)

Architectures

The Figure shows a comparison of the depth and of the performances of the different networks, on the ImageNet challenge.

Revolution of Depth




FIGURE: Evolution of the depth of the CNN and the test error

References

- L. Breiman, J. Friedman, C. J. Stone, R. A. Olshen (1984). *Classification and regression trees*. Chapman et Hall. CRC Press, Boca Raton.
- Giraud C. (2015) *Introduction to High-Dimensional Statistics* Vol. 139 of Monographs on Statistics and Applied Probability. CRC Press, Boca Raton, FL.
- Hastie, T. and Tibshirani, R. and Friedman, J, (2009), *The elements of statistical learning : data mining, inference, and prediction*, Springer.
- R.E. Shapire and Y. Freund *Boosting : Foundation and Algorithms* MIT Press, 2012

References

The [main references](#) for the course on Neural networks are :

- Ian Goodfellow, Yoshua Bengio and Aaron Courville *Deep learning*
- Bishop (1995) *Neural networks for pattern recognition*, Oxford University Press.
- Hugo Larochelle (Sherbrooke) :
<http://www.dmi.usherb.ca/~larocheh/>
- Charles Ollion et Olivier Grisel *Deep learning course*
<https://github.com/m2dsupsdclass/lectures-labs>