

Unsupervised Learning: Clustering

Aurélien Garivier

Formation 2018

Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants

(Agglomerative) Hierarchical Cluster Analysis

Other methods

Clustering

Goal


Goal: automatically discover clusters in the data

- Model-based clustering:
 - assume that the data was generated by a model
 - try to recover the original model from the data
- Model-free clustering:
 - no assumption on the mechanism producing data
 - vector quantization
 - cluster should be *homogeneous* and *different from one another*
 - data-driven loss function to minimize

Model-based clustering: choice of distance

- Data often in \mathbb{R}^P (or projected to)
- Distance sometimes natural, sometimes not
- Often: need to *normalize* first
- Default choice: Euclidian distance $d(x, x') = \sqrt{\sum (x_i - x'_i)^2}$
- Other possible norms: L^1, L^∞ , etc.
- Mahalanobis distance: $d(x, x') = \sqrt{(x - x')\Sigma^{-1}(x - x')}$
- Categorical data: typically χ^2 distance

Ressources

- The Elements of Statistical Learning, *T. Hastie, R. Friedman, J. Tibshirani*, Springer
- Data Mining , *S. Tufféry*, Technip
-  *P. Besse et al.* <http://wikistat.fr/>
- Interesting overview at <http://scikit-learn.org/stable/modules/clustering.html>
- Interesting demos on <https://www.toptal.com/machine-learning/clustering-algorithms>

Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants


(Agglomerative) Hierarchical Cluster Analysis

Other methods

Principal Component Analysis

- Observation X_1, \dots, X_n in \mathbb{R}^p , *centered*
- PCA = Dimension reduction tool (for visual clustering)
- Idea: project onto a subspace $R^d \subset \mathbb{R}^p$ so as to save as much variance as possible
- Big Picture: orthogonal linear transformation such that the first component has highest variance, then the second, etc.
- First component: $w_1 = \operatorname{argmax}_{\|w\|=1} \sum_i (x_i \cdot w)^2$ is the eigenvector of $X^T X$ corresponding to the highest eigenvalue.
- Similar reasoning for the next components in the orthogonal of w_1 .

PCA: visualization



Src: [<https://techannouncer.com/>

global-pca-unit-market-2017-adelte-airmak-industries-amss-ltd-cavotec-airport-division-ciat-effeti/]

PCA algorithm


PCA

- Center all variables
- Compute the $p \times p$ empirical covariance matrix $X^T X$.
- Compute the components $W_d =$ the d first eigenvectors of $X^T X$ in decreasing order of the eigenvalues
- Return the projection of X onto the d first components $T_d = X W_d$.



Then:

- either visualize clusters (2d or 3d plots)
- or use another clustering algorithm on the lower-dimensionnal data T_d (*dimension reduction*)

Example: IRIS



PCA scree plot



Src: <http://strata.uga.edu/8370/lecturenotes/principalComponents.html> and

<http://www.sthda.com/english/wiki/print.php?id=207>

Independent Component Analysis (ICA)

- similar idea, but search for *independent* (instead of uncorrelated) components
- computationally more demanding, but iterative (entropy-based) algorithm exists
- often used for blind source separation
- see also Non-negative Matrix Factorization (NMF)
- for vizualization, see also t-SNE: *t*-distributed Stochastic Neighbor Embedding

Commands

scikitlearn:

```
class sklearn.decomposition.PCA(n_components=None,  
copy=True, whiten=False, svd_solver='auto', tol=0.0,  
iterated_power='auto', random_state=None)
```

ICA: `sklearn.decomposition.FastICA`

R: package stats

```
prcomp() et princomp() [fonction de base, package stats],  
PCA() [package FactoMineR],  
dudi.pca() [package ade4],  
epPCA() [package ExPosition]
```

ICA: `ica`

Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants

(Agglomerative) Hierarchical Cluster Analysis

Other methods

Model-based approach: Gaussian Mixtures

- Observations X_1, \dots, X_n in \mathbb{R}^p
- K cluster centers: μ_1, \dots, μ_K
- Cluster k has probability p_k
- Points in cluster k have law $\mathcal{N}(\mu_k, \Sigma)$
- Given a sample (X_1, \dots, X_n) , how to estimate the cluster centers $(\mu_k)_k$ and how identify the clusters?

Expectation-Maximization Algorithm

- Likelihood: the parameter $\theta = (p, \mu)$ has likelihood

$$L(\theta) \propto \sum_{z_1, \dots, z_n \in \{1, \dots, K\}^n} \prod_{i=1}^n p_{z_i} e^{-\frac{1}{2} \Sigma^{-1} (X_i - \mu_{z_i}) \Sigma^{-1}}$$

\implies non-convex, very hard to maximize

- *Approximate* iterative optimization

Expectation-Maximization Algorithm

Given an estimate $\theta^j = (p^j, \mu^j)$, compute

- *membership weights*

$$w_{i,k}^j = P_{\theta^j}(z_i = k | X_i) = \frac{p_k^j e^{-\frac{1}{2}\Sigma^{-1}(X_i - \mu_k^j)\Sigma^{-1}}}{\sum_{\ell=1}^K p_{\ell}^j e^{-\frac{1}{2}\Sigma^{-1}(X_i - \mu_{\ell}^j)\Sigma^{-1}}}$$

- updated cluster weights:

$$p_k^{j+1} = \frac{\sum_{i=1}^n w_{i,k}^j}{n}$$

- updated cluster means:

$$\mu_k^{j+1} = \frac{\sum_{i=1}^n w_{i,k}^j X_i}{\sum_{i=1}^n w_{i,k}^j}$$

Convergence of the Expectation-Maximization Algorithm

Theorem

The likelihood of the iterates are increasing:

$$L(\theta^{j+1}) \geq L(\theta^j)$$

Good: converges

Bad: local optimum

EM Algorithm

- randomly initialize θ_0
- compute EM iterations until convergence:
 - membership weights $(w_{i,k}^j)_{i,k}$
 - updated cluster weights $(p_k^{j+1})_k$
 - updated cluster means $(\mu_k^{j+1})_k$
- start again (a few times) to look for a better local optimum

Commands

scikitlearn:

```
class sklearn.mixture.GaussianMixture(n_components=1,  
covariance_type='full', tol=0.001, reg_covar=1e-06,  
max_iter=100, n_init=1, init_params='kmeans',  
weights_init=None, means_init=None, precisions_init=None,  
random_state=None, warm_start=False, verbose=0,  
verbose_interval=10)
```

R:

```
MClust() [package MASS],  
GuassianMixtures() [package sBIC]
```

Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants

(Agglomerative) Hierarchical Cluster Analysis

Other methods

Model-free clustering

- Observations X_1, \dots, X_n in \mathbb{R}^p ;
- Objective function: for candidate cluster centers $\mu = (\mu_1, \dots, \mu_K)$ and cluster assignments $z = (z_1, \dots, z_n)$:

$$L(\mu, z) = \sum_{k=1}^K \sum_{i: z_i = k} \|X_i - \mu_k\|^2 = \sum_{i=1}^n \sum_{k=1}^K \mathbb{1}\{z_i = k\} \|X_i - \mu_k\|^2$$

- If $S_k = \{i : z_i = k\}$,

$$L(\mu, z) = \sum_{k=1}^K |S_k| \text{Var}[S_k]$$

- Minimizing L is equivalent to minimizing pairwise deviations in the clusters:

$$\operatorname{argmin}_{\mu, z} L(\mu, z) = \operatorname{argmin}_{\mu, z} \sum_{k=1}^K \frac{1}{|S_k|} \sum_{i, j \in S_k} \|X_i - X_j\|^2$$

Lloyd's algorithm

- For a fixed μ , optimizing in z is easy: choose $z_i = \operatorname{argmin}_k \|X_i - \mu_k\|$
- BUT optimizing in μ is NP-hard!

k-means


- randomly initialize θ_0
- compute Lloyd's iterations until convergence:
 - *membership variables* $z_i^j = \operatorname{argmin}\{\|X_i - \mu_k^j\|$
 - *updated cluster weights* $N_k^j = \sum_{i=1}^n \mathbb{1}\{z_i^j = k\}$
 - *updated cluster means* $\mu_k^{j+1} = \frac{\sum_{i: z_i^j = k} X_i}{N_k^j}$
- start again (a few times) to look for a better local optimum

Comments on k-means

- k-means is a "hard" version of EM for mixtures! (when variances tends to 0)
- Complexity: linear in n and k (fast)
- Requires only a dissimilarity measure (not necessarily a distance)
- Quality of solution found depends on (random) initialization
- Not robust to outliers

- Problem: how to choose k ?

[Src: <http://www.sthda.com/english/articles/29-cluster-validation-essentials/>]



Improved Initialization: k-means++

[Arthur, D.; Vassilvitskii, S. (2007). "k-means++: the advantages of careful seeding" (PDF). Proceedings of the eighteenth annual ACM-SIAM symposium on Discrete algorithms.]

Idea: enforce distant cluster centers from the start

k-means++

- Choose first center μ_1^1 at random
- for $j = 2$ to K , repeat:
 - compute $D_i^j = \min_{\ell < j} \|X_i - \mu_\ell^1\|$
 - choose $\mu_j^1 = X_i$ with probability proportionnal to D_i^j
- theoretical guarantee of $O(\log k)$ approximation from the start
- still linear complexity
- often a dramatic improvement in practice

k-medoids: robustness vs computation time

k-medoids

- randomly initialize θ_0 with K points of the dataset
 - iterate until convergence:
 - *membership variables* $z_i^j = \operatorname{argmin}\{\|X_i - \mu_k^j\|\}$
 - *updated cluster medoids* $\mu_k^{j+1} = \operatorname{argmin}_{i \in S_k} \sum_{\ell \in S_k} \|X_\ell - X_i\|$
 - start again (a few times) to look for a better local optimum
-
- Robust to outliers (cf median versus mean)
 - BUT computation time is quadratic in n
 - k-means is k-medoids with $\|X_\ell - X_i\|^2$ instead of $\|X_\ell - X_i\|$

Commands

scikitlearn:

```
class sklearn.cluster.KMeans(  
n_clusters=8, init='k-means++', n_init=10, max_iter=300,  
tol=0.0001, precompute_distances='auto', verbose=0,  
random_state=None, copy_x=True, n_jobs=1, algorithm='auto'
```

R: package stats

```
kmeans(x, centers, iter.max = 10, nstart = 1,  
algorithm = c("Hartigan-Wong", "Lloyd", "Forgy",  
"MacQueen"), trace=FALSE)
```

Variants to be found in various packages: ClusterR, kmed, etc.

Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants

(Agglomerative) Hierarchical Cluster Analysis

Other methods

Agglomerative clustering

- greedy bottom-up algorithm
- requires a distance (idssimilarity) between observations
 $\|x - x'\|$
- choice of *distance between clusters*:
 - complete linkage: $d(A, B) = \max \{ \|x - x'\| : x \in A, x' \in B \}$
 - single linkage: $d(A, B) = \min \{ \|x - x'\| : x \in A, x' \in B \}$
 - average linkage distance: $d(A, B) = \frac{1}{|A||B|} \sum_{x \in A} \sum_{x' \in B} \|x - x'\|$
 - Ward distance for Euclidian mean:
$$d(A, B) = \frac{|A||B|}{n(|A| + |B|)} \|\bar{A} - \bar{B}\|$$
 - sum of intra-cluster variance
 - etc.

HCA algorithm


HCA

- Initialization: all observations are clusters $\{X_1\}, \dots, \{X_n\}$
- As long as there are at least two clusters:
 - add a link between two clusters with smallest distance
 - merge them for the next iterations
- Return the *dendrogram* = hierarchy of clusters

Property of Ward for Euclidian distance: interclass variance decreasing with the number of classes

Dendrogram

Clustered Iris data set
(the labels give the true flower species)



Author: Talgalili

[https://commons.wikimedia.org/wiki/File:](https://commons.wikimedia.org/wiki/File:Iris_dendrogram.png)

[Iris_dendrogram.png](https://commons.wikimedia.org/wiki/File:Iris_dendrogram.png)

Pros and Cons

- No need to specify the number of clusters in advance
- A relevant choice can be deduced from the observation of the dendrogram (and practical needs)
- Computational complexity in $O(n^2)$
- Does not find an optimal solution

Commands

scikitlearn:

```
class sklearn.cluster.AgglomerativeClustering(n_clusters=2,  
affinity='euclidean', memory=None, connectivity=None,  
compute_full_tree='auto', linkage='ward',  
pooling_func=<function mean>)
```

R:

```
hclust(d, method = "complete", members = NULL)
```


Outline

Introduction

Principal Component Analysis

Model-based clustering: EM algorithm for Gaussian Mixtures

k-means, k-medoids and variants

(Agglomerative) Hierarchical Cluster Analysis

Other methods

Affinity Propagation

- Message-passing algorithm
- affinity function $s(x, x')$, for example $s(x, x') = -\|x - x'\|^2$
- $s(x, x) =$ input preference : the lower, the higher the chances to be an exemplar
- responsibility matrix R : $r(i, k) =$ how well-suited X_k can serve as an exemplar for X_i (wrt other candidate exemplars)
- availability matrix A : $a(i, k) =$ how appropriate it is for X_i to pick X_k as an exemplar, taking into account other points preferences for X_k as an exemplar

Affinity Propagation

Affinity Propagation

- Initialize R and A with 0
- Repeat until convergence:

- update responsibilities:

$$r(i, k) \leftarrow s(i, k) - \max_{k' \neq k} a(i, k') + s(i, k')$$

- update availabilities:

$$a(i, k) \leftarrow \min \left\{ 0, r(k, k) + \sum_{i' \notin \{i, k\}} \max(0, r(i', k)) \right\} \text{ for } i \neq k$$
$$\text{and } a(k, k) \leftarrow \sum_{i' \neq k} \max(0, r(i', k))$$

- Pick exemplars as maximizers of $r(i, i) + a(i, i)$

Affinity Propagation: pros/cons

- No need to specify the number of clusters
- ... but a parameter plays the same role
- quadratic time complexity
- some improvement over k-means in some cases

Commands

scikitlearn:

```
class sklearn.cluster.AffinityPropagation(damping=0.5,  
max_iter=200, convergence_iter=15, copy=True,  
preference=None, affinity='euclidean', verbose=False)
```

R: package APCluster

```
apcluster(s, x, p=NA, q=NA, maxits=1000,  
convits=100, lam=0.9, includeSim=FALSE, details=FALSE,  
nonoise=FALSE, seed=NA)
```

Spectral Clustering

- Similarity matrix $S_{i,j}$, for example $S_{i,j} = -\|X_i - X_j\|^2$
- Idea: use standard clustering method on eigenvectors of the (normalized) Laplacian matrix

$$L = Id - D^{-1/2} S D^{-1/2}$$

where D is diagonal with $D_{i,i} = \sum_j S_{i,j}$

- Intuition: if S is diagonal by blocks, the eigenvectors are the indicators of the blocks

Commands

scikitlearn:

```
class sklearn.cluster.SpectralClustering(n_clusters=8,  
eigen_solver=None, random_state=None, n_init=10, gamma=1.0,  
affinity='rbf', n_neighbors=10, eigen_tol=0.0,  
assign_labels='kmeans', degree=3, coef0=1,  
kernel_params=None, n_jobs=1)
```

R: package kernlab

```
specc(x, data = NULL, na.action = na.omit, ...)
```

DBscan

- DBscan = density-based spatial clustering of applications with noise
- parameters: radius ϵ and minimal cluster size *minSize*

DBscan

Repeat as long as at least one point has not been visited:

- pick an unvisited point X_i at random
- if it has less than *minSize* ϵ -neighbors, mark it as outlier
- other, form the cluster of all points that can be reached by jumps of at most ϵ starting from X_i

DBscan: comments

- simple and fast
- no need to specify number of clusters
- Problem: sensitive to the choice of parameters
- choosing the right parameters ϵ and *minSize* properly is hard
- choice of ϵ : such that the proportion of outliers is at most 10% (say)
- choice of *minSize*: such that at least 90% have at least *minSize* neighbors
- unable to handle clusters with very different densities

Commands

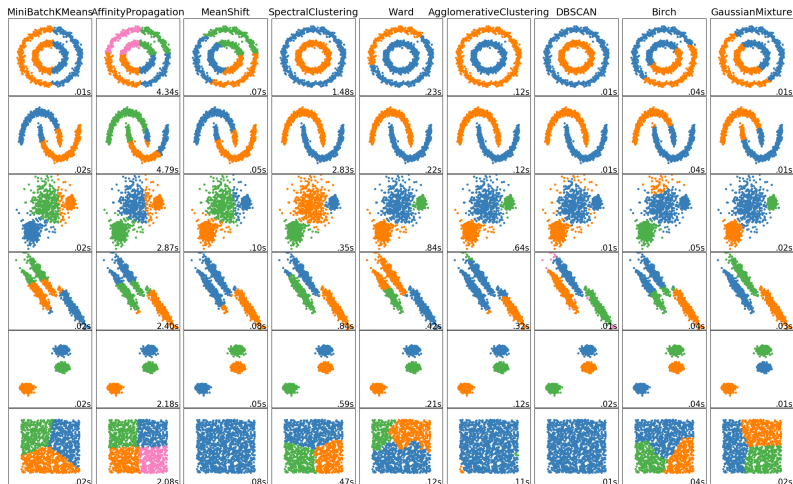
scikitlearn:

```
class sklearn.cluster.DBSCAN(eps=0.5, min_samples=5,  
metric='euclidean', metric_params=None, algorithm='auto',  
leaf_size=30, p=None, n_jobs=1)
```

R: package dbscan

```
dbscan(x, eps, minPts = 5, weights = NULL,  
borderPoints = TRUE, ...)
```

Which algorithm to choose?



Src: [http://scikit-learn.org/stable/auto_examples/cluster/plot_cluster_comparison.html]

Clustering time series

Features:

- mean
- trend
- auto-correlation coefficients
- inter-series correlation
- etc.

⇒ it depends on the nature of the problem and on the goal of the clustering!